

AI6104 - MATHEMATICS FOR AI

TUTORIAL 6 - MATRIX CALCULUS

Problem 1

Let A be an invertible $m \times m$ matrix whose elements are functions of a scalar parameter α . Prove that

$$\frac{\partial A^{-1}}{\partial \alpha} = -A^{-1} \frac{\partial A}{\partial \alpha} A^{-1}$$

Hint: by the definition of inversion, we have $A^{-1}A = I$.

Problem 2

We will look at the backpropagation in one layer of a neural network (see Figure 1).

In neural networks, a layer f is a function of input X and weight W , where the output is $Y = f(X, W)$. In this problem, we assume a linear layer, $Y = f(X, W) = XW$. If we consider the input X has N samples, and each sample, $x^{(i)}$, is a D -dimensional vector, then X is an $N \times D$ matrix. Similarly, we have corresponding N output vectors, and each $y^{(i)}$ is M -dimensional, which forms a $N \times M$ matrix Y . Thus, the weight matrix W has shape $D \times M$. Similar layers as described above are embedded into larger neural networks with loss, usually a scalar, L .

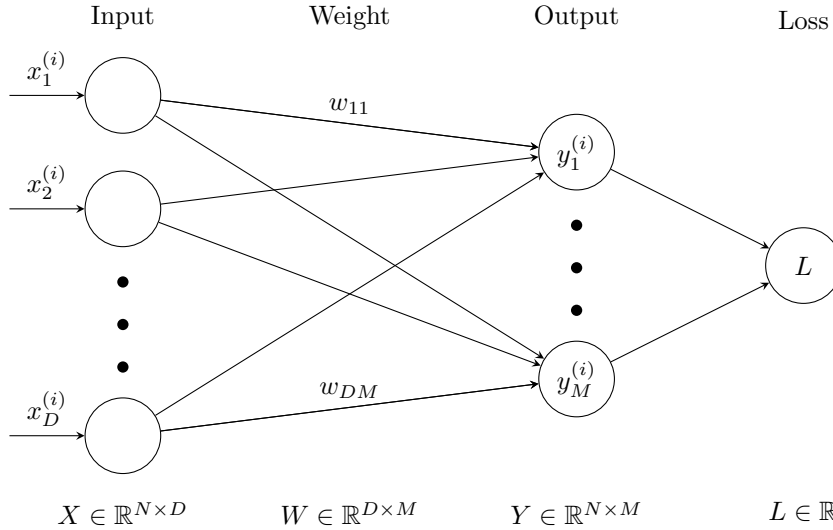


Figure 1: An illustration for Problem 2

- During backpropagation, we want to know the gradient of loss with respect to the input X and weight W . Write down $\frac{\partial L}{\partial X}$ and $\frac{\partial L}{\partial W}$ using chain rule.
- According to part (a), discuss the difficulties of calculating $\frac{\partial L}{\partial X}$ and $\frac{\partial L}{\partial W}$ using the matrix multiplications explicitly.

Problem 3

Suppose a linear layer, $Y = f(X, W)$, in a neural network has activation function $A(Y)$. Calculate $\frac{\partial A}{\partial Y}$ for the following different activation functions.

- (a) **ReLU**: $R(Y) = \max\{0, Y\}$
- (b) **Sigmoid**: $S(Y) = \frac{1}{1+e^{-Y}}$
- (c) **Tanh**: $\tanh(Y) = \frac{e^Y - e^{-Y}}{e^Y + e^{-Y}}$

Problem 4

This problem requires computing the gradients of a full neural network. In particular we are going to compute the gradients of a neural network with one hidden layer trained with cross-entropy loss. The forward pass of single D -dimensional input sample x is as follows:

$$\begin{aligned}
 x &= \text{input} \in \mathbb{R}^D \\
 z &= Wx \in \mathbb{R}^M \\
 h &= \text{ReLU}(z) \in \mathbb{R}^M \\
 \theta &= Uh \in \mathbb{R}^C \\
 \hat{y} &= \text{softmax}(\theta) \in \mathbb{R}^C \\
 J &= \text{CrossEntropy}(y, \hat{y}) \in \mathbb{R}
 \end{aligned}$$

where W and U are weight matrices. y is the true label vector. The softmax activation of the j -th output unit is

$$\hat{y}_j = \text{softmax}(\theta_j) = \frac{e^{\theta_j}}{\sum_j e^{\theta_j}}$$

The cross-entropy error function is

$$\text{CrossEntropy}(y, \hat{y}) = - \sum_j y_j \log \hat{y}_j$$

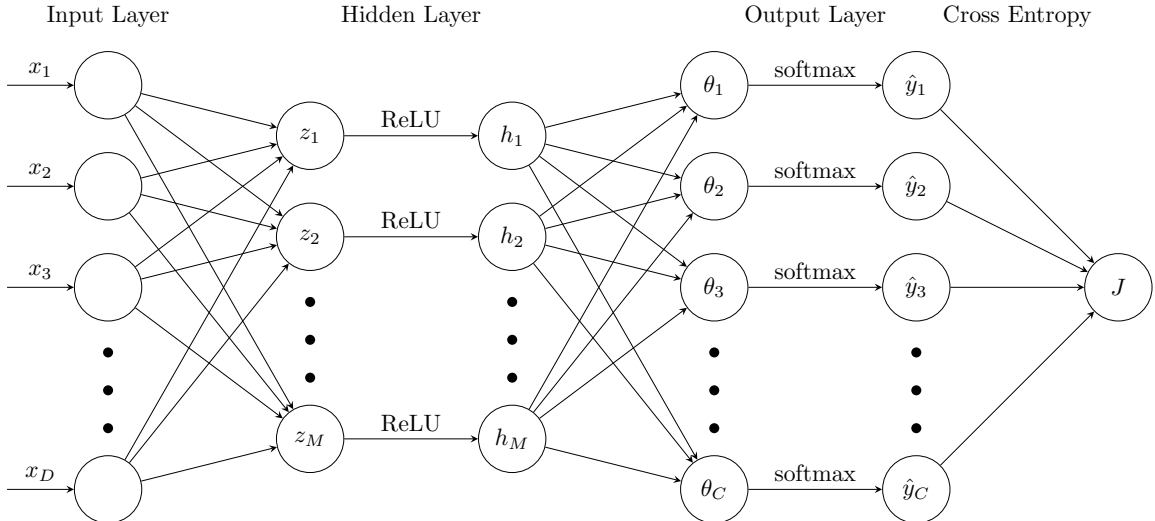


Figure 2: A neural network with one hidden layer

- (a) Calculate the gradient of cross entropy error with respect to the logits, *i.e.*, $\frac{\partial J}{\partial \theta}$.
- (b) According to above information, compute all of the network's gradients, $\frac{\partial J}{\partial U}$, $\frac{\partial J}{\partial W}$, $\frac{\partial J}{\partial x}$.
Hint: In practice, you may transpose your vectors to match the dimensions.